

SQL 1 : Basic Statements

Gary KL Tam

Department of Computer Science
Swansea University

- Structured query language (**SQL**) is a user-friendly language for specifying relational algebra queries. It is supported by all the major database systems.

SQL provides:

- Data Manipulation Language (DML)
 - retrieve, insert and modify database contents
- Data Definition Language (DDL)
 - add and delete database objects
- Data Control Language (DCL)
 - configure security access
- In this lecture, we will learn how to rewrite algebra operators in SQL (DML).

Syntax of an SQL Statement

```
select distinct  $A_1, A_2, \dots, A_n$   
from  $T_1, \dots, T_m$   
where  $P$ 
```

where T_1, \dots, T_m are tables, A_1, \dots, A_n are attributes, and P is a predicate. The statement returns a table, and corresponds to the following relational algebra query:

$$\Pi_{A_1, \dots, A_n}(\sigma_P(T_1 \times \dots \times T_m))$$

Selection σ

```
select *  
from  $T$   
where  $P$ 
```

corresponds to

$$\sigma_P(T)$$

PROF

pid	name	dept	rank	sal
<i>p1</i>	Adam	CS	asst	6000
<i>p2</i>	Bob	EE	asso	8000
<i>p3</i>	Calvin	CS	full	10000
<i>p4</i>	Dorothy	EE	asst	5000
<i>p5</i>	Emily	EE	asso	8500
<i>p6</i>	Frank	CS	full	9000

```
select *
from PROF
where rank = 'asst'
```

$$\sigma_{\text{rank} = \text{"asst"}}(\text{PROF})$$

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

```
select *
from PROF
where not(rank = 'asst' and dept = 'EE')
```

$\sigma_{\neg(\text{rank} = \text{"asst"} \wedge \text{dept} = \text{"EE"})}(\text{PROF})$

Selection Predicate

```
select *  
from  $T$   
where  $P$ 
```

In P , you can specify the standard comparisons and logic operators:

- $=, <>, <, <=, >, >=$
- Connect multiple comparisons with: AND, OR, NOT.

Projection Π

```
select distinct  $A_1, \dots, A_n$   
from  $T$ 
```

corresponds to

$$\Pi_{A_1, \dots, A_n}(T)$$

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

select distinct dept, rank
from PROF

$\Pi_{\text{dept, rank}}(\text{PROF})$

Note

The keyword **distinct** removes all duplicate rows in the output. Omitting the keyword keeps all duplicates. See the next slide.

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

“select dept, rank from PROF” returns:

dept	rank
CS	asst
EE	asso
CS	full
EE	asst
EE	asso
CS	full

This duplicate-retaining feature is useful for aggregate queries as we will discuss later in the course.

Cartesian Product \times

select *
from T_1, T_2

corresponds to $T_1 \times T_2$

select *
from T_1, \dots, T_m

corresponds to $T_1 \times \dots \times T_m$

PROF

pid	name	dept	rank	sal
<i>p1</i>	Adam	CS	asst	6000
<i>p2</i>	Bob	EE	asso	8000
<i>p3</i>	Calvin	CS	full	10000
<i>p4</i>	Dorothy	EE	asst	5000
<i>p5</i>	Emily	EE	asso	8500

TEACH

pid	cid	year
<i>p1</i>	<i>c1</i>	2011
<i>p2</i>	<i>c2</i>	2012
<i>p1</i>	<i>c2</i>	2012

```
select *
from PROF, TEACH
```

PROF × TEACH

Putting Multiple Operators Together

PROF

pid	name	dept	rank	sal
<i>p1</i>	Adam	CS	asst	6000
<i>p2</i>	Bob	EE	asso	8000
<i>p3</i>	Calvin	CS	full	10000
<i>p4</i>	Dorothy	EE	asst	5000
<i>p5</i>	Emily	EE	asso	8500

TEACH

pid	cid	year
<i>p1</i>	<i>c1</i>	2011
<i>p2</i>	<i>c2</i>	2012
<i>p1</i>	<i>c2</i>	2012

select distinct dept
from PROF, TEACH
where PROF.pid = TEACH.pid

$\Pi_{\text{dept}}(\sigma_{\text{PROF.pid} = \text{TEACH.pid}}(\text{PROF} \times \text{TEACH}))$

Rename ρ

```
select ...  
from  $T$  as  $S$   
where ...
```

corresponds to

$$\dots \rho_S(T) \dots$$

PROF

pid	name	dept	rank	sal
<i>p1</i>	Adam	CS	asst	6000
<i>p2</i>	Bob	EE	asso	8000
<i>p3</i>	Calvin	CS	full	10000
<i>p4</i>	Dorothy	EE	asst	5000
<i>p5</i>	Emily	EE	asso	8500

TEACH

pid	cid	year
<i>p1</i>	<i>c1</i>	2011
<i>p2</i>	<i>c2</i>	2012
<i>p1</i>	<i>c2</i>	2012

select distinct dept
 from PROF as A, TEACH as B
 where A.pid = B.pid

$$\Pi_{\text{dept}}(\sigma_{A.\text{pid} = B.\text{pid}}(\rho_A(\text{PROF}) \times \rho_B(\text{TEACH})))$$

([SQL statement 1])
minus
([SQL statement 2])

corresponds to

$$T_1 - T_2$$

where T_1 (T_2) is the table returned by SQL statement 1 (2).

Note

- T_1 and T_2 need to have the same schema.
- Duplicates in T_1 and T_2 will first be removed before performing the set difference.
- In some systems (e.g., SQL server from Microsoft), the set difference operator is named “except”, instead of “minus”.

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

(select rank from PROF)

minus

(select rank from PROF where dept = 'CS')

$$\Pi_{\text{rank}}(\text{PROF}) - \Pi_{\text{rank}}(\sigma_{\text{dept} = \text{"CS"}}(\text{PROF}))$$

([SQL statement 1])
union
([SQL statement 2])

corresponds to

$$T_1 \cup T_2$$

where T_1 (T_2) is the table returned by SQL statement 1 (2).

Note

- T_1 and T_2 need to have the same schema.
- Duplicates in T_1 and T_2 will first be removed before performing the set union.

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

```
(select * from PROF where sal <= 6000)
union
(select * from PROF where sal >= 9000)
```

$$\sigma_{sal \leq 6000}(\text{PROF}) \cup \sigma_{sal \geq 9000}(\text{PROF})$$

We have shown how to rewrite the 6 fundamental algebra operators in SQL. How about the extended operators \leftarrow , \cap , \bowtie and \div ? As we will see next, there is an explicit statement only for \cap . Nevertheless, as \cap and \bowtie can be implemented using the 6 fundamental operators, they can also be written in SQL using the statements introduced earlier. We will, however, ignore \leftarrow from our discussion (this operator is the least useful one, anyway).

Set Intersection \cap

([SQL statement 1])
intersect
([SQL statement 2])

corresponds to

$$T_1 \cap T_2$$

where T_1 (T_2) is the table returned by SQL statement 1 (2).

Note

- T_1 and T_2 need to have the same schema.
- Duplicates in T_1 and T_2 will first be removed before performing the set union.

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

(select * from PROF where sal >= 6000)
 intersect
 (select * from PROF where dept = 'CS')

$$\sigma_{sal \geq 6000}(\text{PROF}) \cap \sigma_{dept = \text{"CS"}}(\text{PROF})$$

Natural Join

PROF

pid	name	dept	rank	sal
<i>p1</i>	Adam	CS	asst	6000
<i>p2</i>	Bob	EE	asso	8000
<i>p3</i>	Calvin	CS	full	10000
<i>p4</i>	Dorothy	EE	asst	5000
<i>p5</i>	Emily	EE	asso	8500

TEACH

pid	cid	year
<i>p1</i>	<i>c1</i>	2011
<i>p2</i>	<i>c2</i>	2012
<i>p1</i>	<i>c2</i>	2012

select distinct PROF.pid, name, dept, rank, sal, cid, year
from PROF, TEACH
where PROF.pid = TEACH.pid

$\Pi_{\text{PROF.pid, name, dept, rank, sal, cid, year}}(\sigma_{\text{PROF.pid=TEACH.pid}}(\text{PROF} \times \text{TEACH}))$

=

$\text{PROF} \bowtie \text{TEACH}$

Division

T_1		T_2	
pid	cid	cid	
p1	c1	c1	
p1	c2	c2	
p1	c3	c3	
p2	c2		
p2	c3		
p3	c1		
p4	c1		
p4	c2		
p4	c3		

(select pid from T_1)
minus
select pid from (
select * from (select pid from T_1), T_2)
minus
(select * from T_1))

Note

Notice how an SQL statement can be nested in a from clause.

$$\Pi_{S_1-S_2}(T_1) - \Pi_{S_1-S_2}(\Pi_{S_1-S_2}(T_1) \times T_2 - T_1) = T_1 \div T_2$$

Funny stuffs

```
SELECT * FROM politicians WHERE clue > 0;
```

```
SQL> select intelligence_level from developer;  
select intelligence_level from developer  
      *  
ERROR at line 1:  
ORA-00904: "INTELLIGENCE_LEVEL": invalid identifier
```

http://www.orafaq.com/wiki/Fun_stuff