

Week 9

Sorting

- 1 Binary heaps
- 2 Heapification
- 3 Building a heap
- 4 HEAP-SORT
- 5 Priority queues
- 6 QUICK-SORT
- 7 Analysing QUICK-SORT
- 8 Exercises
- 9 Outlook

Binary heaps

Heapification

Building a
heap

HEAP-
SORT

Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook

- We return to sorting, considering HEAP-SORT and QUICK-SORT.

Reading from CLRS for week 7

- 1 Chapter 6, Sections 6.1 - 6.5.
- 2 Chapter 7, Sections 7.1, 7.2.

Discover the properties of binary heaps

Running example

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT



Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook



First property: level-completeness

In your first year you have seen **binary trees**:

- ❶ In data structures they should be as “balanced” as possible, i.e., their height should be as small as possible for the given number of nodes.
- ❷ Recall that the height $\text{ht}(T)$ of a rooted tree is the maximum distance between the root and any leaf.
- ❸ Perfect are the perfect binary trees (for height h they have $2^{h+1} - 1$ nodes).
- ❹ Now close to perfect come the **level-complete binary trees**:
 - ❶ We can partition the nodes of a (binary) tree T into levels, according to their distance from the root.
 - ❷ We have levels $0, 1, \dots, \text{ht}(T)$.
 - ❸ In general, level k has from 1 to 2^k nodes.
 - ❹ If all levels k except possibly for the last one are full (have precisely 2^k nodes in them), then we call the tree **level-complete**.

Binary heaps

Heapification


Building a
heapHEAP-
SORTPriority
queuesQUICK-
SORTAnalysing
QUICK-
SORT

Exercises


Outlook

Examples

The binary tree



is level-complete (level-sizes are 1, 2, 4, 4), while



is not (level-sizes are 1, 2, 3, 6).

The height of level-complete binary trees

For a level-complete binary tree T we have

$$\text{ht}(T) = \lfloor \lg(\# \text{nds}(T)) \rfloor,$$

where $\# \text{nds}(T)$ is the number of nodes of T .

That is, the height of T is the binary logarithm of the number of nodes of T , after removal of the fractional part.

- “Balanced” (“good”) trees T should have

$$\text{ht}(T) \approx \lg(\# \text{nds}(T)).$$

- Level-complete binary trees definitely fulfil that.

Operations in binary trees are often
linear in the height of the tree.

Thus the height should be as small as possible.

Binary heaps

Heapification

Building a
heapHEAP-
SORTPriority
queuesQUICK-
SORTAnalysing
QUICK-
SORT

Exercises

Outlook

Second property: completeness

To have simple and efficient access to the nodes of the tree, the nodes of the last layer better are not placed in random order:

- Best is if they fill the positions from the left without gaps.
- A level-complete binary tree with such gap-less last layer is called a **complete tree**.
- So the level-complete binary tree on the examples-slide is not complete.
- While the running-example is complete.

Third property: the heap-property

The running-example is NOT a binary *search* tree:

- 1 It would be too expensive to have this property (all nodes to the left \leq , all nodes to the right \geq) *together* with the completeness property.
- 2 However we have another property related to order (not just related to the structure of the tree): The value of every node is not less than the value of any of its successors (the nodes below it).
- 3 In other words, all nodes below (to the left *and* to the right) are \leq .
- 4 This property is called the **heap property**.
- 5 More precisely it is the **max-heap property**.

Definition 1

A **binary heap** is a binary tree which is complete and has the heap property. More precisely we have **binary max-heaps** and **binary min-heaps**.

[Binary heaps](#)[Heapification](#)[Building a heap](#)[HEAP-SORT](#)[Priority queues](#)[QUICK-SORT](#)[Analysing QUICK-SORT](#)[Exercises](#)[Outlook](#)

Fourth property: Efficient index computation

Consider the numbering (not the values) of the nodes of the running-example:

- 1 This numbering follows the layers, beginning with the first layer and going from left to right.
- 2 Due to the completeness property (no gaps!) these numbers yield easy relations between a parent and its children.
- 3 If the node has number p , then the left child has number $2p$, and the right child has number $2p + 1$.
- 4 And the parent has number $\lfloor p/2 \rfloor$.

Efficient array implementation

For binary search trees we needed full-fledged trees (as you implemented in your first year):

- 1 That is, we needed nodes with three pointers: to the parent and to the two children.
- 2 However now, for complete binary trees we can use a more efficient array implementation, using the numbering for the array-indices.

So a binary heap with m nodes is represented by an array with m elements:

- C-based languages use 0-based indices (while the book uses 1-based indices).
- For such an index $0 \leq i < m$ the index of the left child is $2i + 1$, and the index of the right child is $2i + 2$.
- While the index of the parent is $\lfloor (i - 1)/2 \rfloor$.

Binary heaps




Heapification

Building a
heapHEAP-
SORTPriority
queuesQUICK-
SORTAnalysing
QUICK-
SORT

Exercises

Outlook

Float down a single disturbance



The idea of heapification

- The input is an array A and index i into A .
- It is assumed that the binary trees rooted at the left and right child of i are binary (max-)heaps, but we do not assume anything on $A[i]$.
- After the “heapification”, the values of the binary tree rooted at i have been rearranged, so that it is a binary (max-)heap now.

For that, the algorithm proceeds as follows:

- 1 First the largest of $A[i]$, $A[l]$, $A[r]$ is determined, where $l = 2i$ and $r = 2i + 1$ (the two children).
- 2 If $A[i]$ is largest, then we are done.
- 3 Otherwise $A[i]$ is swapped with the largest element, and we call the procedure recursively on the changed subtree.

Analysing heapification

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT

Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook

Obviously, we go down from the node to a leaf (in the worst case), and thus the running-time of heapification is







linear in the height h of the subtree.

This is $O(\lg n)$, where n is the number of nodes in the subtree (due to $h = \lfloor \lg n \rfloor$).

Heapify bottom-up

A

4	1	3	2	16	9	10	14	8	7
---	---	---	---	----	---	----	----	---	---



The idea of building a binary heap

One starts with an arbitrary array A of length n , which shall be re-arranged into a binary heap. Our example is

$$A = (4, 1, 3, 2, 16, 9, 10, 14, 8, 7).$$

We repair (heapify) the binary trees bottom-up:

- 1 The leaves (the final part, from $\lfloor n/2 \rfloor + 1$ to n) are already binary heaps on their own.
- 2 For the other nodes, from right to left, we just call the heapify-procedure.

Roughly we have $O(n \cdot \lg n)$ many operations:

- 1 Here however it pays off to take into account that most of the subtrees are small.
- 2 Then we get run-time $O(n)$.

So building a heap is linear in the number of elements.

Heapify and remove from last to first

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT











Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook



A

1	2	3	4	7	8	9	10	14	16
---	---	---	---	---	---	---	----	----	----

(k)

The idea of HEAP-SORT

Now the task is to sort an array A of length n :

- ❶ First make a binary max-heap out of A (in linear time).
- ❷ Repeat the following until $n = 1$:
 - ❶ The maximum element is now $A[1]$ — swap that with the last element $A[n]$, and remove that last element, i.e., set $n := n - 1$.
 - ❷ Now perform heapification for the root, i.e., $i = 1$. We have a binary (max-)heap again (of length one less).

The run-time is $O(n \cdot \lg n)$.

Remark on SELECTION-SORT

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT

Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook

We see that HEAP-SORT is an intelligent form of SELECTION-SORT (where we move from right to left, not from left to right):

- 1 We extract the largest element, the second largest etc., as does SELECTION-SORT.
- 2 But we support the search by the binary-heap data structure.
- 3 Thus each search takes only time $O(\lg n)$, where n is the number of remaining elements, instead of $\Theta(n)$ as for SELECTION-SORT.

All basic operations are (nearly) there

Recall that a (basic) (max-)priority queue has the operations:

- MAXIMUM
- DELETE-MAX
- INSERTION.

We use an array A containing a binary (max-)heap (the task is just to maintain the max-heap-property!):

- 1 The maximum is $A[1]$.
- 2 For deleting the maximum element, we put the last element $A[n]$ into $A[1]$, decrease the length by one (i.e., $n := n - 1$), and heapify the root (i.e., $i = 1$).
- 3 And we add a new element by adding it to the end of the current array, and heapifying all its predecessors up on the way to the root.

Using our running-example, a few slides ago for HEAP-SORT:

- 1 Considering it from (a) to (j), we can see what happens when we perform a sequence of DELETE-MAX operations, until the heap only contains one element (we ignore here the shaded elements — they are visible only for the HEAP-SORT).
- 2 And considering the sequence in reverse order, we can see what happens when we call INSERTION on the respective first shaded elements (these are special insertions, always inserting a new max-element).

- MAXIMUM is a constant-time operation.
- DELETE-MAX is one application of heapification, and so need time $O(\lg n)$ (where n is the current number of elements in the heap).
- INSERTION seems to invoke up to the current height many applications of heapification, and thus would look like $O((\lg n)^2)$, but it's easy to see that it is $O(\lg n)$ as well (see the tutorial).

The idea of QUICK-SORT

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT

Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook

Remember MERGE-SORT:

- A divide-and-conquer algorithm for sorting an array in time $O(n \cdot \lg n)$.
- The array is split in half, the two parts are sorted recursively (via MERGE-SORT), and then the two sorted half-arrays are merged to the sorted (full-)array.

Now we split along an element x of the array:

- We partition into elements $\leq x$ (first array) and $> x$ (second array).
- Then we sort the two sub-arrays recursively.
- Done!

Remark on ranges

In the book arrays are 1-based:

- ① So the indices for an array A of length n are $1, \dots, n$.
- ② Accordingly, a sub-array is given by indices $p \leq r$, meaning the range p, \dots, r .

For Java-code (C, C++, C#, etc.) we use 0-based arrays:

- ① So the indices are $0, \dots, n - 1$.
- ② Accordingly, a sub-array is given by indices $p < r$, meaning the range $p, \dots, r - 1$.

Range-bounds for a sub-array are here now always
left-closed and right-open!

So the whole array is given by the range-parameters $0, n$.

The idea of partitioning in-place

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT


Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook



An example



Instead of i we use $q = i + 1$:

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ≡ ≡ ≡ ↺ 🔍 ↻

Selecting the pivot


The partitioning-procedure expects the partitioning-element to be the last array-element. So for selecting the pivot, we can just choose the last element:

```
private static void  
    place_partition_element_last(final int []  
        A, final int p, final int r) {}
```

However this makes it vulnerable to “malicious” choices, so we better **randomise**:

```
private static void  
    place_partition_element_last(final int []  
        A, final int p, final int r) {  
    final int i = p+(int) Math.random()*(r-p);  
    {final int t=A[i]; A[i]=A[r-1]; A[r-1]=t;}  
}
```

A not unreasonable tree



Average-case

If we actually achieve that both sub-arrays are at least a constant fraction α of the whole array (in the previous picture, that's $\alpha = 0.1$), then we get

$$T(n) = T(\alpha \cdot n) + T((1 - \alpha) \cdot n) + \Theta(n).$$

That's basically the second case of the Master Theorem (the picture says it's similar to $\alpha = \frac{1}{2}$), and so we *would* get

$$T(n) = \Theta(n \cdot \log n).$$

And we actually get that:

- for the non-randomised version (choosing always the last element as pivot), when averaging over all possible input sequences (without repetitions);
- for the randomised version (choosing a random pivot), when averaging over all (internal!) random choices; here we do not have to assume something on the inputs, except that all values are different.

However, as the tutorial shows:

The worst-case run-time of QUICK-SORT is $\Theta(n^2)$
(for both versions)!

- This can be repaired, making also the worst-case run-time $\Theta(n \cdot \log n)$.
- For example by using median-computation in linear time for the choice of the pivot.
- In practice these days often (for example in gcc) Introsort is used, which starts with QUICK-SORT, and switches to HEAP-SORT if recursion depth becomes too big, while when the array is small enough, INSERTION-SORT is used.
- When pointers are to be handled, then Timsort is preferable, which is derived from MERGE-SORT and INSERTION-SORT.

HEAP-SORT on sorted sequence

What does HEAP-SORT on an already sorted sequence? And what's the complexity? Consider the input sequence

$1, 2, \dots, 10.$

HEAP-SORT on sorted sequence (cont.)

Now the sorting works as follows:

- ❶ 10, 9, 7, 8, 5, 6, 3, 1, 4, 2
- ❷ 9, 8, 7, 4, 5, 6, 3, 1, 2; 10
- ❸ 8, 5, 7, 4, 2, 6, 3, 1; 9, 10
- ❹ 7, 5, 6, 4, 2, 1, 3; 8, 9, 10
- ❺ 6, 5, 3, 4, 2, 1; 7, 8, 9, 10
- ❻ 5, 4, 3, 1, 2; 6, 7, 8, 9, 10
- ❼ 4, 2, 3, 1; 5, 6, 7, 8, 9, 10
- ❽ 3, 2, 1; 4, 5, 6, 7, 8, 9, 10
- ❾ 2, 1; 3, 4, 5, 6, 7, 8, 9, 10
- ❿ 1; 2, 3, 4, 5, 6, 7, 8, 9, 10

So HEAP-SORT doesn't take any advantage from some already existing sorting.

[Binary heaps](#)[Heapification](#)[Building a heap](#)[HEAP-SORT](#)[Priority queues](#)[QUICK-SORT](#)[Analysing QUICK-SORT](#)[Exercises](#)[Outlook](#)

Simplifying insertion

When discussing insertion into a (max-)priority-queue, implemented via a binary (max-)heap, we just used a general addition of one element into an existing heap:

- The insertion-procedure used heapification up on the path to the root.
- Now actually we have always special cases of heapification — namely which?

Simplifying insertion

When discussing insertion into a (max-)priority-queue, implemented via a binary (max-)heap, we just used a general addition of one element into an existing heap:

- The insertion-procedure used heapification up on the path to the root.
- Now actually we have always special cases of heapification — namely which?

Answer: The new element only *goes up*, swapping with its parent, which then has already its place, until the new element found its place.

Change to the partitioning procedure

What happens if we change the line

```
if (v <= x) {A[j] = A[q]; A[q++] = v;}
```

of function `partition` to

```
if (v < x) {A[j] = A[q]; A[q++] = v;}
```

- Can we do it?
- Would it have advantages?

Change to the partitioning procedure

What happens if we change the line

```
if (v <= x) {A[j] = A[q]; A[q++] = v;}
```

of function partition to

```
if (v < x) {A[j] = A[q]; A[q++] = v;}
```

- Can we do it?
- Would it have advantages?

Answer: Yes, we can do it: it would save work, due to values equal x not being moved to the left part.

QUICK-SORT on constant sequences

What is QUICK-SORT doing on a constant sequence, in its three incarnations:

- pivot is last element
- pivot is random element
- pivot is median element?

QUICK-SORT on constant sequences

What is QUICK-SORT doing on a constant sequence, in its three incarnations:

- pivot is last element
- pivot is random element
- pivot is median element?

One of the two sub-arrays will have size 1, and QUICK-SORT degenerates to an $O(n^2)$ algorithm (which does nothing).

What can we do about it?

QUICK-SORT on constant sequences

What is QUICK-SORT doing on a constant sequence, in its three incarnations:

- pivot is last element
- pivot is random element
- pivot is median element?

One of the two sub-arrays will have size 1, and QUICK-SORT degenerates to an $O(n^2)$ algorithm (which does nothing).

What can we do about it?

We can refine the partition-procedure by

- not just splitting into two parts,
- but into three parts: all elements $< x$, all elements $= x$, and all elements $> x$.

Then we use only the left and the right sub-array (leaving out the whole middle) for recursion. We get $O(n \log n)$ for constant sequences.

Worst-case for QUICK-SORT

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT

Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook

Consider sequences without repetitions, and assume the pivot is always the last element:

- What is a worst-case input?
- And what is QUICK-SORT doing on it?

Worst-case for QUICK-SORT

Consider sequences without repetitions, and assume the pivot is always the last element:

- What is a worst-case input?
- And what is QUICK-SORT doing on it?

Every already sorted sequence is a worst-case example!
QUICK-SORT behaves as with constant sequences.

Note that this is avoided with randomised pivot-choice (and, of course, with median pivot-choice).

Worst-case $O(n \log n)$ for QUICK-SORT

How can we achieve $O(n \log n)$ in the **worst-case** for QUICK-SORT?

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT

Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook

Worst-case $O(n \log n)$ for QUICK-SORT

How can we achieve $O(n \log n)$ in the **worst-case** for QUICK-SORT?

- Just choosing, within our current framework, the median-*element* is not enough, but we need the change the framework, allowing to compute the median-*index* (or we do the triple-partition as discussed).
- Best is to remove the function `place_partition_element_last`, and leave the partitioning fully to function `partition`.

Then the main procedure becomes (without the asserts):

```
public static void sort(final int [] A, final
    int p, final int r) {
    if (r-p <= 1) return;
    final int q = partition(A,p,r);
    sort(A,p,q); sort(A,q+1,r);
}
```

Divide-and-Conquer

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT

Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook

- Matrix Multiplication is a great example, which you should have seen.
- Otherwise, if you do not have special interests, I think you can move on to other paradigms.
- Sure, algorithmic analysis (and the underlying mathematics) is an infinite game, but you might learn it as you go.
- The Wikipedia page on the Master Theorem is quite nice.

- Graph theory has many many aspects.
- Much theory, much experimentation.
- “Network” are graphs, but with a specialised focus on growth of the graph.

Data structures

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT

Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook

- The book has plenty on the classical data structures.
- Hashing is a very active topic, with special links to cryptography.

Complexity theory

CS.270
Algorithms

Oliver
Kullmann

Binary heaps

Heapification

Building a
heap

HEAP-
SORT

Priority
queues

QUICK-
SORT

Analysing
QUICK-
SORT

Exercises

Outlook

- This is about about the resource needs of algorithmic problem IN GENERAL.
- Most central (and most famous) the “P versus NP” question: is there a polytime algorithm for finding an assignment for a propositional formula?

- SAT is the field of deciding whether propositional formulas are satisfiable (have an assignment making them true), or not (are unsatisfiable).
- See The Science of Brute Force.